



Your agent has memory. Does it have a *brain* to remember from?

How I run two companies and my whole life on one AI agent, three Obsidian "walls," and a hundred skills. And how you can wire the same idea into your Hermes.

Renuka • co-founder, PraxisLoop + Raygency • scroll to begin

— THE PREMISE

Memory is not the same as a second brain.

Hermes, G-Brain, Mem0, Honcho: these give your agent *memory of your conversations*. That is real, and it matters.

But a question they do not answer: where does your actual **work**, your **knowledge**, and your **many-context life** physically live? An agent that remembers what you said, but has no organized place to read your world from, is recalling a pile, not navigating a library.

This is the layer I built first. Everything that follows is how.

— MEET THE TEAM

Two residents run this house.

★ RESIDENT CELEBRITY



Evie // my Claude Code agent

The tabby with the stripes, and the name behind my local agent. Evie reads and writes my Obsidian vaults, runs my skills, holds my memory, and keeps me on top of my days. She is the operator of the second brain.

★ RESIDENT CELEBRITY



Enzo // my Hermes agent

White with the grey spots (front of frame). Enzo is my Hermes setup, living on my co-founder Varun's Mac. A separate environment, a separate brain, and the agent this whole talk is really about helping you build.

— THE SHAPE OF IT

One operating system. Two companies. One life.

1

agent (Evie) at the center

3

Obsidian walls, fully isolated

100⁺

skills built, a rotation used weekly

144

memory files, loaded every session

None of this is theory. It runs every day. Standups, brain-dumps, content for seven brands, outreach, and the supervision of a second agent all flow through it.

What a Tuesday actually looks like.

07:00

Morning standup

Evie reads my FOCUS file and hands me the day's three priorities and any fixed-time anchor. `/morning-standup`

all day

I brain-dump on my phone

Every stray thought goes into a voice note. No filtering, no organizing while talking. `Pocket`

evening

Evie sorts the pile

She pulls the recordings, extracts tasks and decisions, and routes each one to the correct wall. `/pocket`

on cue

The work runs itself

Weekly content for each brand, replies to outreach, competitor ads turned into briefs. `/weekly-content` · `/route-outreach-reply` · `/reverse-engineer-ad`

close

She remembers so I do not have to

Session-end updates memory and writes the handoff, so tomorrow's Evie wakes up knowing today. `/session-end`

Three vaults, not one. On purpose.

A single vault would let two companies bleed into each other. Separate walls give hard isolation, the right co-founder sees only their world, and every skill is fenced to read only its own wall.

WALL 01

PraxisLoop

Consulting and curriculum. Clients, prospects, engagements, research outputs.

synced with → Greg

WALL 02

Raygency

The AI product studio. Products, agents, UGC, speaking, go-to-market.

synced with → Varun

WALL 03

Evie (personal)

Cross-company and personal. Knowledge, journals, health, infra, the day's focus.

private → just me + Evie

Three walls. Three terminals.

Each wall gets its own live agent session, open side by side, each fenced to one company. On a normal day I am running three Evies at once.

● ● ● PRAXISLOOP

PL ▶ signal scan + Greg sync
█

● ● ● RAYGENCY

RG ▶ kreashot content + Enzo
█

● ● ● EVIE

EV ▶ inbox, focus, memory
█

~/ .claude/learnings/today.md Each terminal is an isolated context window: they cannot read each other. So every one drops a trace into **one shared daily ledger**, and at night **/daily** reads it to reconstruct the whole day across all three.

The same shape, every single time.

Raygency/

- ├ 00-Inbox // raw capture lands here
- ├ 01-Clients
- ├ 02-Prospects
- ├ 04-Knowledge // the durable stuff
- ├ 05-Content
- ├ 09-Meetings // call notes, by person
- ├ 12-Products
- └ _Claude-Bridge.md // the index Evie reads

Because the taxonomy is identical across all three walls, Evie never has to guess where something goes. A meeting note always lands in **09-Meetings**. A captured idea always starts in **00-Inbox**.

Predictable structure is what turns "store this somewhere" into automatic routing. The repetition is the feature.

— HOW THE AGENT CONNECTS

The mechanism is boringly simple.

A vault is just a folder of markdown files on your machine.

So the agent needs exactly one thing: permission to read and write that folder. No special API, no server, no cloud. Evie runs as Claude Code with direct file access, and the vault is right there on disk.

read

write

search

link

That local-file simplicity is also the privacy model. The vaults never leave my machine. Sync to my co-founders is handled by Obsidian itself, end to end, not by any agent.

Proven, not claimed: this is exactly how Evie works today. The vault-to-agent wiring on the Claude Code side is the part I have actually run for months. Hold that thought for the Hermes translation.

I built over a hundred. I do not use them all.

A skill is a repeatable workflow the agent invokes by name. Building one is cheap; the discipline is knowing which ones earn a place in the weekly rotation. Here is what mine actually is, right now.

DAILY RHYTHM

/morning-standup

/pocket

/session-end

/log

CONTENT ENGINE

/weekly-content ×7

/x-autopost

EVIE ↔ ENZO OPS

/enzo-review

/evie-call

/strategic-lenses

INBOX & COMMS

/inbox-triage

/draft-reply

/capture

OUTREACH

/route-outreach-reply

/signal-engine-pl

/reverse-engineer-ad

ON A SCHEDULE · CRON

morning-brew · 7am

newsletter · Thu

podcast · Thu

enzo-watch · 3×/day

"How did you make that ad?"

You saw the Kreashot cut. Here is the honest answer: it was built on my own machine with a code-driven video pipeline, orchestrated by Evie. No Hermes, no no-code tool.



01 Brief

A competitor ad reverse-engineered into a structured beat sheet (hook, arc, CTA) as data. [/reverse-engineer-ad](#)

02 Render engine

My RayMotion repo: a Remotion (React) template that renders any ad from that brief, seven beat layouts. [BriefAd](#) • [Remotion v4](#)

03 Voiceover

A HeyGen avatar (Maya) speaks the script; ffmpeg pulls clean audio out of the render. [HeyGen](#) • [ffmpeg](#)

04 Moving hero + B-roll

Real footage plus generative clips where needed, dropped into the editorial beats. [Higgsfield Seedance](#) • [Pexels](#)

05 Quality gate

An automated check on brightness and saturation, with a branded frame-0 poster as the cover. [qa-gate.sh](#)

THE STACK

[Remotion v4](#) [React / TypeScript](#) [HeyGen](#) [Higgsfield](#) [Pexels](#) [Whisper](#) [ffmpeg](#) [Claude Code \(Evie\)](#)

The point for this room: code-driven video means every render is repeatable and editable in seconds. Evie drives it. It is the local sibling of the agent setup, not a separate magic tool.

144 files. One index. Loaded every session.

Evie's memory is not a black box. It is human-readable files plus a master index pulled into context at the start of every session. It maps cleanly onto the five-layer model Vignesh teaches.

1

Working memory

the live session plus the handoff note

→ session + handoff.md

2

Episodic

growth log, moods, what worked, what did not

→ feedback + growth_log

3

Semantic

durable facts in topic files and the knowledge base

→ memory/ + 04-Knowledge

4

Cross-session retrieval

the master index loaded at every start

→ MEMORY.md

5

Skill memory

reusable workflows and standard procedures

→ .claude/skills/

Evie is teaching Enzo to be better.

I gave Evie access to the Telegram group where Enzo, my Hermes agent, lives. Because Evie is the more mature agent, she watches what Enzo is doing and sends him instructions on my behalf, guiding his work while I sleep or step away.

One agent supervising and cross-training another. The structured second brain on one side raising the conversational agent on the other.

Evie reads Enzo's Telegram

guides his builds

reviews his output

What's next: giving Enzo his own Obsidian vault too, so I can see just how powerful a Hermes agent becomes once it has a real second brain behind it.



Evie (back) keeping an eye on Enzo (front). As it should be.

How to wire this into your own agent.

My setup is proven on Claude Code. The *principle* ports straight to Hermes; the exact wiring is the piece each of you verifies on your own machine. Three moves:

1

Give Hermes the folder

Install Obsidian, create your vault, and give your Hermes agent file access to that vault directory on the same machine. Same-box filesystem access is the whole trick.

2

Mirror the taxonomy

Build one repeated folder structure across every wall, and teach it to Hermes once as a rule. Predictable shape is what makes routing automatic.

3

Add skills per context

Wrap your real workflows as named skills, each fenced to one wall. Start with capture and a daily standup. Grow only what you use.

Honest line for the room: I have not yet wired Obsidian into Enzo myself. I have run this for months on Claude Code, and giving Enzo his own vault is my next build. So treat steps as the architecture to translate, and verify the Hermes specifics as you go. No hand-waving.

The honest FAQ.

Do I need to know how to code?

No. The agent reads and writes the files for you. Even the video pipeline is code the agent writes; you direct it in plain language.

Does it work offline?

The vault and files are fully local. The model (the agent's brain) needs internet, unless you self-host one, which needs serious hardware.

Do I still need G-Brain if I have this?

Yes. Different jobs: G-Brain remembers conversations, Obsidian holds your knowledge. Keep both.

Does the agent ever get it wrong?

Yes, so I supervise. It drafts, I approve. Nothing important ships on the agent's say-so alone.

Mac or Windows?

Either. Obsidian and Hermes both run on both. The vault is just a folder on whichever machine your agent lives on.

Free or paid?

Obsidian is free. I pay for **Sync** so I can share walls with my co-founders. Your model is the real cost.

How do I keep two businesses from mixing?

Separate vaults (walls), and every skill fenced to read only its own wall. Hard isolation by design.

Where do I start?

One vault, the repeated folder structure, and two skills: **capture** and a **daily standup**. Grow only what you use.

"If I have G-Brain, do I even need Obsidian?"

G-BRAIN · MEMO · HONCHO

Agent-facing memory

A database the agent retrieves from. Structured for the machine. You do not browse it. It answers: "what does my agent remember about me?"

OBSIDIAN · THE WALLS

Human-facing brain

Files you and the agent both read, link, and navigate, split by context. It answers: "where does my actual knowledge and work live?"

They are partners, not replacements. Keep both.



Thank you. Now go build yours.

Renuka, Evie & Enzo